

Query Processing and Optimization

Week - 3 ,Topic – 2, Lesson 5 and 6



Outline

- ❑ Query optimization techniques (rule-based + practical)
- ❑ Cost-based optimization and execution plans
- ❑ Use case with examples

Pipelining (Streaming)

- ❑ **Iterator model:** each operator produces the next tuple on demand (`next()`).
- ❑ **Pipelining:** pass tuples directly from one operator to the next (no full intermediate table).
 - ❑ Pipelines can be executed in two ways: demand driven and producer driven
- ❑ **Materialization:** write intermediate results to disk/temp table (slower, but sometimes required).
- ❑ Pipeline breaks: sort, hash build phase, or blocking aggregates often need full input.

Example pipeline:

Scan → Filter (σ) → Project (π) → Output

`SELECT Name FROM Student WHERE Marks > 80;`

Implementation of demand-driven pipelining

Each operation is implemented as an **iterator** implementing the following operations

open()

- file scan: initialize file scan
- state: pointer to beginning of file
- merge join: sort relations;
- state: pointers to beginning of sorted relations

next()

- for file scan: Output next tuple, and advance and store file pointer
- for merge join: continue with merge from earlier state till next output tuple is found. Save pointers as iterator state.

close()

3: What is Query Optimization?

3. Query optimization techniques

- ❑ Optimization chooses a good plan among many equivalent plans.
- ❑ Logical optimization: rewrite query without changing meaning (algebraic rules).
- ❑ Physical optimization: pick access paths + join algorithms + join order.
- ❑ Two styles: rule-based (heuristics) and cost-based (use a cost model).

Classic Heuristic Rules

Simple rules that often help

- Push selections down: apply σ as early as possible (reduce rows early).
- Push projections down: keep only needed columns (reduce row width).
- Combine selections: $\sigma_{c_1}(\sigma_{c_2}(R)) \rightarrow \sigma_{c_2} \wedge_{c_2}(R)$
- Reorder joins (when safe): join smaller / filtered relations first.
- Replace Cartesian Products: Use joins with conditions instead of cross products.
- Query Tree: a graphical representation of the operators, relations, attributes and predicates and processing sequence during query processing. It is composed of three main parts:
 - The Leafs: the base relations used for processing the query/ extracting the required information.
 - The Root: the final result/relation as an output based on the operation on the relations used for query processing
 - Nodes: intermediate results or relations before reaching the final result.
 - Sequence of execution of operation in a query tree will start from the leaves and continues to the intermediate nodes and ends at the root.

Example Cont...

Example-

Before optimization


```
SELECT E.name FROM Employee E, Department D  
WHERE E.dept_id = D.dept_id AND D.location = 'Binh Doung';
```

Problem


- Cartesian product produces huge intermediate results
- Selection is applied too late

After Optimization

```
SELECT E.name FROM Employee E JOIN Department D ON  
E.dept_id = D.dept_id WHERE D.location = ' Binh Doung ';
```



Query tree –



Another Example (Push Selection)

Same result, usually faster



TRƯỜNG ĐẠI HỌC
QUỐC TẾ
MIỀN ĐÔNG
EASTERN
INTERNATIONAL
UNIVERSITY

Original (conceptual):

$R \bowtie S$ then $\sigma_{\{R.x > 10\}}$

Rewrite:

$(\sigma_{\{R.x > 10\}}(R)) \bowtie S$

- Filtering R first shrinks the join input.
- Less data enters the join → fewer comparisons and less I/O.

Physical Aids (Indexes & Views)

Optimization uses what exists in storage



TRƯỜNG ĐẠI HỌC
QUỐC TẾ
MIỀN ĐÔNG
EASTERN
INTERNATIONAL
UNIVERSITY

- Indexes can turn scans into fast lookups (B+ tree, hash index).
- Clustered index: table stored in index order (good for ranges).
- Materialized view: precomputed result stored on disk.
- Partitioning: split table to scan only relevant parts.

Optimizer chooses plans based on available indexes/views + estimated benefit.

Statistics & Selectivity (Basics)

Estimates guide the optimizer



TRƯỜNG ĐẠI HỌC
QUỐC TẾ
MIỀN ĐÔNG
EASTERN
INTERNATIONAL
UNIVERSITY

- DBMS stores stats: #rows, #pages, distinct values, histograms.
- Selectivity estimates how many rows pass a predicate.
- Common assumptions: uniform distribution and independence (not always true).
- Better stats → better plan choices (especially for join order).

Example: if attribute A has V distinct values, $\text{sel}(A = c) \approx 1 / V$ (uniform).

4: Cost-Based Optimization (CBO)

4. Cost-based optimization and execution plans



TRƯỜNG ĐẠI HỌC
QUỐC TẾ
MIỀN ĐÔNG
EASTERN
INTERNATIONAL
UNIVERSITY

- CBO picks the plan with the lowest estimated cost.
- Cost is usually dominated by disk I/O, then CPU, Network cost (in distributed databases), then memory.
- DBMS uses table/index statistics to estimate: cardinalities and costs.
- The chosen plan is an execution plan (physical operator tree).
- Highly efficient and Adapts to data size and distribution

Techniques Used

- Evaluates multiple query plans
- Uses statistics about tables and indexes

Chooses optimal:

- Join order
- Join method (Nested Loop, Hash Join, Merge Join)
- Access path (Index scan vs. Table scan)

A Simple Cost Model

Not exact—just useful for comparisons

- I/O Cost : Represents the cost of reading and writing data pages from disk. Since disk access is slow, I/O cost is usually the largest contributor to total cost.
 - I/O cost = Number of page reads / writes × Cost per I/O operation
- CPU Cost: Represents the processing cost of evaluating tuples (rows). Includes comparisons, joins, filtering, and aggregation.
 - CPU cost = Number of tuples processed × Cost per tuple operation
- Memory Cost (Impact): Memory is not always counted directly.

If sufficient memory is available:

- ✓ Operations like sorting and hashing stay in memory → **lower I/O cost**

If memory is insufficient:

- ✓ Intermediate results are written to disk → **higher I/O cost**

The total cost of an execution plan is often estimated using:

$$\text{Cost(plan)} = (N_{IO} \times t_{IO}) + (N_{CPU} \times t_{CPU})$$

Where

N_{IO} = Number of page I/O operations

t_{IO} = Time per I/O operation

N_{CPU} = Number of tuple (row) operations

t_{CPU} = Time per CPU operation

Cost Based Optimization Example

Assume

Employee has 1,000 rows

Department has 10 rows

Index exists on Employee.dept_id

Query -

```
SELECT name, dept_name FROM Employee E,  
Department D WHERE E.dept_id = D.dept_id;
```

Possible Execution Plans

Plan 1: Nested Loop Join (Employee → Department)

- Outer table: Employee (1,000 rows)
- Inner table: Department (10 rows)
- Cost $\approx 1,000 \times 10 = 10,000$ comparisons

Plan 2: Index Join (Department → Employee)

- Outer table: Department (10 rows)
- Use index on Employee.dept_id
- Cost ≈ 10 index lookups = very low

Optimizer chooses Plan 2 because it has lower estimated cost based on:

- Table size
- Index availability
- Number of comparisons



Name	Dept_name
Kim	IT
Nygen	IT
Tien	HR

More Example....

Assume the database system uses the following costs:

Time per I/O operation

$$t_{IO} = 5 \text{ ms}$$

Time per CPU operation

$$t_{CPU} = 0.01$$

Plan A: Full Table Scan

Number of page I/Os

$$N_{IO} = 200$$

Number of tuple operations

$$N_{CPU} = 10000$$

$$\text{Sol: Cost(A)} = (200 \times 5) + (10,000 \times 0.01)$$

$$= 1,000 + 100$$

$$= 1,100 \text{ ms}$$

Plan B: Index Scan

Number of page I/Os

$$N_{IO} = 50$$

Number of tuple operations

$$N_{CPU} = 2,000$$

Sol:

$$\text{Cost(B)} = (50 \times 5) + (2,000 \times 0.01)$$

$$= 250 + 20$$

$$= 270 \text{ ms}$$

emp_id	name
101	Kim
105	Nyugen

Searching Join Orders (Idea Only)

Enumerating plans without exploding



- Join order matters a lot ($A \bowtie B \bowtie C$ may be much faster than $A \bowtie (B \bowtie C)$).
- Optimizers often search “left-deep” trees (one join at a time).
- Dynamic programming idea: reuse best plan for each subset of tables.
- Heuristics/pruning stop the search from becoming too large.

Key point: the optimizer explores alternatives and keeps the cheapest (estimated) ones.

Reading an Execution Plan

What you see in EXPLAIN output



- Operators: Seq Scan, Index Scan, Hash Join, Sort, Aggregate, etc.
- Estimated rows/cardinality at each node.
- Estimated cost (often “startup..total”).
- Actual runtime (if you run EXPLAIN ANALYZE) may differ from estimates.

Mini example (conceptual):

Hash Join

- > Seq Scan on R
- > Hash
- > Index Scan on S

Employee Use Case Example -

```
CREATE TABLE Employee (emp_id INT, name  
VARCHAR(20), dept_id INT, salary INT);
```

```
CREATE TABLE Department (dept_id INT,  
dept_name VARCHAR(20), location VARCHAR(20));
```

```
INSERT INTO Employee VALUES (1, 'Kim', 10, 60000), (2,  
'Nyugen', 20, 45000), (3, 'Nahi', 10, 70000), (4, 'Tien', 30,  
40000);
```

```
INSERT INTO Department VALUES (10, 'IT', 'Binh Dung'),  
(20, 'HR', 'Ho Chi Minh City'), (30, 'Sales', 'Hanoi');
```

Classical JOIN Query:

```
SELECT E.name, D.dept_name FROM Employee E,  
Department D WHERE E.dept_id = D.dept_id AND  
D.location = 'Binh Dung' AND E.salary > 50000;
```

Heuristic (Rule-Based) Optimization - Optimized SQL (Logical)

Rule Applied - Combine selections, Push selection down, Replace Cartesian product with JOIN and Project required columns early

```
SELECT E.name, D.dept_name FROM Employee E  
JOIN Department D ON E.dept_id = D.dept_id  
WHERE D.location = 'Binh Dung' AND E.salary >  
50000;
```

Employee Use Case Example -

Cost-Based Optimization:

Assume information-

- Employee table = 1000 rows (200 pages)
- Department table = 50 rows (5 pages)
- Index on Employee.dept_id
- Disk I/O cost $t_{IO} = 5 \text{ ms}$
- CPU cost $t_{CPU} = 0.01 \text{ ms}$

Plan B: Index Join (Chosen Plan)

$$N_{IO} = 55 \text{ pages}$$

$$N_{CPU} = 2,000 \text{ tuple operations}$$

$$\begin{aligned} \text{Cost}(B) &= (55 \times 5) + (2,000 \times 0.01) \\ &= 275 + 20 \\ &= 295 \text{ ms} \end{aligned}$$

Cost-based optimizer selects Plan B


Plan A: Full Table Scan + Nested Loop Join

Solutions-

$$N_{IO} = 205 \text{ pages}$$

$$N_{CPU} = 10,000 \text{ tuple operations}$$

$$\begin{aligned} \text{Cost}(A) &= (205 \times 5) + (10,000 \times 0.01) \\ &= 1025 + 100 \\ &= 1125 \text{ ms} \end{aligned}$$



Revision

- Query processing: parse → rewrite → optimize → execute.
- Joins have multiple algorithms; pick based on size, indexes, and predicates.
- Optimization uses rules + statistics to reduce work early.
- Cost-based optimizers compare many plans and choose the cheapest estimate.

Tip for practice: write a query, guess the plan, then check with EXPLAIN in a DBMS.

Further Reading and Practice

- * Silberschatz, A., Korth, H. F., & Sudarshan, S. (2020). Database System Concepts (7th ed.). McGraw-Hill. (Part six, Chapter 16. Query optimization page no-743 to 794)
- * Rupley Jr, Michael L. "Introduction to query processing and optimization." Indiana University at South Bend (2008).
- * Practice Relational database with SQL compiler Such as MySQL, PostgreSQL, SQL Workbench
 - * Online compiler - <https://sqlfiddle.com/>
https://sqlzoo.net/wiki/SQL_Tutorial